

Reinforcement Learning


Solving MDPs

0.51	0.72	0.84	1.00
0.27		0.55	-1.00
0.00	0.22	0.37	0.13

VALUES AFTER 5 ITERATIONS

Marcello Restelli

March–April, 2015





Brute Force

Marcello
Restelli

Policy Search

Dynamic
Programming

Policy Iteration

Value Iteration

Extensions to

Dynamic
Programming

Linear

Programming

- Solving an MDP means finding an **optimal policy**
- A **naive** approach consists of
 - **enumerating** all the deterministic Markov policies
 - **evaluate** each policy
 - **return** the best one
- The number of policies is **exponential**: $|\mathcal{A}|^{|S|}$
- Need a **more intelligent search** for best policies
 - **restrict the search** to a subset of the possible policies
 - using **stochastic optimization** algorithms



What is Dynamic Programming?

Marcello
Restelli

Policy Search

Dynamic
Programming

Policy Iteration

Value Iteration

Extensions to

Dynamic

Programming

Linear

Programming

- **Dynamic**: sequential or temporal component to the problem
- **Programming**: optimizing a “program”, i.e., a policy
 - c.f. linear programming
- A method for solving **complex** problems
- By breaking them down into **subproblems**
 - **Solve** the subproblems
 - **Combine** solutions to subproblems



Requirements for Dynamic Programming

Marcello
Restelli

Policy Search

Dynamic
Programming

Policy Iteration

Value Iteration

Extensions to

Dynamic

Programming

Linear
Programming

- Dynamic Programming is a **very general** solution method for problems which have **two properties**:
 - **Optimal substructure**
 - **Principle of optimality** applies
 - Optimal solution can be decomposed into **subproblems**
 - **Overlapping subproblems**
 - Subproblems **recur** many times
 - Solutions can be **cached** and **reused**
- Markov decision processes satisfy both properties
 - **Bellman equation** gives recursive decomposition
 - **Value function** stores and reuses solutions



Planning by Dynamic Programming

Marcello
Restelli

Policy Search

Dynamic
Programming

Policy Iteration

Value Iteration

Extensions to

Dynamic

Programming

Linear

Programming

- Dynamic Programming assumes **full knowledge** of the MDP
- It is used for **planning** in an MDP
- **Prediction**
 - Input: MDP $\langle \mathcal{S}, \mathcal{A}, P, R, \gamma, \mu \rangle$ and policy π (i.e., MRP $\langle \mathcal{S}, P^\pi, R^\pi, \gamma, \mu \rangle$)
 - Output: value function V^π
- **Control**
 - Input: MDP $\langle \mathcal{S}, \mathcal{A}, P, R, \gamma, \mu \rangle$
 - Output: value function V^* and optimal policy π^*



Other Applications of Dynamic Programming

Marcello
Restelli

Policy Search

Dynamic
Programming

Policy Iteration

Value Iteration

Extensions to

Dynamic

Programming

Linear

Programming

Dynamic Programming is used to solve many other problems:

- Scheduling algorithms
- String algorithms (e.g., sequence alignment)
- Graph algorithms (e.g., shortest path algorithms)
- Graphical models (e.g., Viterbi algorithm)
- Bioinformatics (e.g., lattice models)



Finite–Horizon Dynamic Programming

Marcello
Restelli

Policy Search

Dynamic
Programming

Policy Iteration

Value Iteration

Extensions to

Dynamic
Programming

Linear
Programming

- **Principle of optimality**: the tail of an optimal policy is optimal for the “tail” problem
- **Backward induction**
 - **Backward recursion**

$$V_k^*(s) = \max_{a \in \mathcal{A}_k} \left\{ R_k(s, a) + \sum_{s' \in \mathcal{S}_{k+1}} P_k(s' | s, a) V_{k+1}^*(s') \right\}, \quad k = N - 1, \dots, 0$$

- **Optimal policy**

$$\pi_k^*(s) \in \arg \max_{a \in \mathcal{A}_k} \left\{ R_k(s, a) + \sum_{s' \in \mathcal{S}_{k+1}} P_k(s' | s, a) V_{k+1}^*(s') \right\}, \quad k = 0, \dots, N - 1$$

- **Cost**: $N|\mathcal{S}||\mathcal{A}|$ vs $|\mathcal{A}|^{N|\mathcal{S}|}$ of brute force policy search
- From now on, we will consider **infinite–horizon discounted** MDPs



Policy Evaluation

Marcello
Restelli

Policy Search

Dynamic
Programming

Policy Iteration
Value Iteration
Extensions to
Dynamic
Programming

Linear
Programming

- For a **given policy** π compute the **state-value function** V^π
- Recall
 - State-value function for policy π :

$$V^\pi(s) = \mathbb{E} \left\{ \sum_{t=0}^{\infty} \gamma^t r_t \mid s_0 = s \right\}$$

- **Bellman equation** for V^π :

$$V^\pi(s) = \sum_{a \in \mathcal{A}} \pi(a|s) \left[R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V^\pi(s') \right]$$

- A **system** of $|\mathcal{S}|$ simultaneous **linear equations**
- Solution in **matrix** notation (complexity $O(n^3)$):

$$V^\pi = (I - \gamma P^\pi)^{-1} R^\pi$$



Iterative Policy Evaluation

Marcello
Restelli

Policy Search

Dynamic
Programming

Policy Iteration
Value Iteration

Extensions to
Dynamic
Programming

Linear
Programming

- Iterative application of Bellman expectation backup
- $V_0 \rightarrow V_1 \rightarrow \dots \rightarrow V_k \rightarrow V_{k+1} \rightarrow \dots \rightarrow V^\pi$
- A **full policy–evaluation backup**:

$$V_{k+1}(s) \leftarrow \sum_{a \in \mathcal{A}} \pi(a|s) \left[R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V_k(s') \right]$$

- A **sweep** consists of applying a backup operation to each state
- Using **synchronous** backups
 - At each iteration $k + 1$
 - For all states $s \in \mathcal{S}$
 - Update $V_{k+1}(s)$ from $V_k(s')$



Example

Small Gridworld

Marcello
Restelli

Policy Search

Dynamic
Programming

Policy Iteration

Value Iteration


Extensions to

Dynamic

Programming

Linear

Programming



actions

	1	2	3
4	5	6	7
8	9	10	11
12	13	14	

$r = -1$
on all transitions

- **Undiscounted episodic MDP**
 - $\gamma = 1$
 - All episodes terminate in **absorbing** terminal state
- **Transient** states $1, \dots, 14$
- One **terminal** state (shown twice as shaded squares)
- Actions that would take agent off the grid leave state **unchanged**
- Reward is -1 until the terminal state is reached



Policy Evaluation in Small Gridworld

Marcello
Restelli

Policy Search

Dynamic
Programming

Policy Iteration

Value Iteration

Extensions to

Dynamic

Programming

Linear
Programming

V_k for the
Random Policy

$k = 0$

0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0

$k = 1$

0.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	0.0

$k = 2$

0.0	-1.7	-2.0	-2.0
-1.7	-2.0	-2.0	-2.0
-2.0	-2.0	-2.0	-1.7
-2.0	-2.0	-1.7	0.0

Greedy Policy
w.r.t. V_k

	↔	↔	↔
↔	↔	↔	↔
↔	↔	↔	↔
↔	↔	↔	

← random
policy

	←	↔	↔
↑	↔	↔	↔
↔	↔	↔	↓
↔	↔	→	

	←	←	↔
↑	↔	↔	↓
↑	↔	↔	↓
↔	→	→	



Policy Evaluation in Small Gridworld

Marcello
Restelli

Policy Search

Dynamic
Programming

Policy Iteration

Value Iteration

Extensions to

Dynamic


Programming

Linear

Programming


$k = 3$

0.0	-2.4	-2.9	-3.0
-2.4	-2.9	-3.0	-2.9
-2.9	-3.0	-2.9	-2.4
-3.0	-2.9	-2.4	0.0




$k = 10$

0.0	-6.1	-8.4	-9.0
-6.1	-7.7	-8.4	-8.4
-8.4	-8.4	-7.7	-6.1
-9.0	-8.4	-6.1	0.0



$k = \infty$

0.0	-14.	-20.	-22.
-14.	-18.	-20.	-20.
-20.	-20.	-18.	-14.
-22.	-20.	-14.	0.0



optimal
policy



Policy Improvement

Marcello
Restelli

Policy Search

Dynamic
Programming

Policy Iteration

Value Iteration

Extensions to

Dynamic

Programming

Linear

Programming

- Consider a **deterministic policy** π
- For a given state s , would it **better** to do an action $a \neq \pi(s)$?
- We can **improve** the policy by acting greedily

$$\pi'(s) = \arg \max_{a \in \mathcal{A}} Q^\pi(s, a)$$

- This improves the value from **any** state s over one step

$$Q^\pi(s, \pi'(s)) = \max_{a \in \mathcal{A}} Q^\pi(s, a) \geq Q^\pi(s, \pi(s)) = V^\pi(s)$$



Policy Improvement Theorem

Marcello
Restelli

Policy Search

Dynamic
Programming

Policy Iteration

Value Iteration

Extensions to

Dynamic

Programming

Linear

Programming

Theorem

Let π and π' be any pair of deterministic policies such that

$$Q^\pi(s, \pi'(s)) \geq V^\pi(s) \quad , \quad \forall s \in \mathcal{S}$$

Then the policy π' must be as good as, or better than π

$$V^{\pi'}(s) \geq V^\pi(s) \quad , \quad s \in \mathcal{S}$$

Proof.

$$\begin{aligned} V^\pi(s) &\leq Q^\pi(s, \pi'(s)) = \mathbb{E}_{\pi'} [r_{t+1} + \gamma V^\pi(s_{t+1}) | s_t = s] \\ &\leq \mathbb{E}_{\pi'} [r_{t+1} + \gamma Q^\pi(s_{t+1}, \pi'(s_{t+1})) | s_t = s] \\ &\leq \mathbb{E}_{\pi'} [r_{t+1} + \gamma r_{t+2} + \gamma^2 Q^\pi(s_{t+2}, \pi'(s_{t+2})) | s_t = s] \\ &\leq \mathbb{E}_{\pi'} [r_{t+1} + \gamma r_{t+2} + \dots | s_t = s] = V^{\pi'}(s) \end{aligned}$$





Policy Iteration

Marcello
Restelli

Policy Search

Dynamic
Programming

Policy Iteration

Value Iteration

Extensions to

Dynamic
Programming

Linear

Programming

- What if improvements **stops** ($V^{\pi'} = V^{\pi}$)?

$$Q^{\pi}(s, \pi'(s)) = \max_{a \in \mathcal{A}} Q^{\pi}(s, a) = Q^{\pi}(s, \pi(s)) = V^{\pi}(s)$$

- But this is the **Bellman optimality equation**
- Therefore $V^{\pi}(s) = V^{\pi'}(s) = V^*(s)$ for all $s \in \mathcal{S}$
- So π is an **optimal** policy!

$$\pi_0 \rightarrow V^{\pi_0} \rightarrow \pi_1 \rightarrow V^{\pi_1} \rightarrow \dots \rightarrow \pi^* \rightarrow V^* \rightarrow \pi^*$$



Example of Policy Iteration

Jack's Car Rental

Marcello
Restelli

Policy Search

Dynamic
Programming

Policy Iteration

Value Iteration

Extensions to
Dynamic
Programming

Linear
Programming

- **States:** Two locations, maximum of 20 cars each
- **Actions:** Move up to 5 cars between two locations overnight
- **Reward:** \$10 for each car rented (must be available)
- **Transitions:** Cars returned and requested randomly
 - **Poisson distribution**, n returns/request with probability $\frac{\lambda^n}{n!} e^{-\lambda}$
 - **First location:** average requests = 3, average returns = 3
 - **Second location:** average requests = 4, average returns = 2



Example of Policy Iteration

Jack's Car Rental

Marcello
Restelli

Policy Search

Dynamic
Programming

Policy Iteration

Value Iteration


Extensions to

Dynamic

Programming

Linear

Programming





Modified Policy Iteration

Marcello
Restelli

Policy Search

Dynamic
Programming

Policy Iteration
Value Iteration

Extensions to
Dynamic
Programming

Linear
Programming

- Does policy evaluation **need to converge** to V^π ?
- Or should we introduce a **stopping condition**
 - e.g., ϵ -convergence of value function
- Or simply **stop after k iterations** of iterative policy evaluation?
- For example, in the small gridworld $k = 3$ was sufficient to achieve optimal policy
- Why not update policy **every iteration**? i.e. stop after $k = 1$



Generalized Policy Iteration

Marcello
Restelli

Policy Search

Dynamic
Programming

Policy Iteration

Value Iteration


Extensions to

Dynamic


Programming

Linear

Programming



- **Policy evaluation:** Estimate V^π
 - e.g., Iterative policy evaluation
- **Policy improvement:** Generate $\pi' \geq \pi$
 - e.g., Greedy policy improvement





Principle of Optimality

Deterministic Shortest Path

Marcello
Restelli

Policy Search

Dynamic
Programming

Policy Iteration

Value Iteration

Extensions to

Dynamic

Programming

Linear

Programming

- **Problem:** reach goal state g from start state s_1 with min cost
- Each action a from state s leads to:
 - **Deterministic** transition $P(s'|s, a) = 1$ iff $s' = \text{succ}(s, a)$
 - Reward $R(s, a)$ negative (is a **cost**)
 - **Undiscounted** $\gamma = 1$

Theorem

A path $s_1, a_1, s_2, a_2, \dots, s_T$ is optimal if and only if:

- For any **intermediate** state s_t along the solution path
- s_t, a_t, \dots, s_T is an **optimal path** from s_t to g



Principle of Optimality

Marcello
Restelli

Policy Search

Dynamic
Programming

Policy Iteration

Value Iteration

Extensions to

Dynamic

Programming

Linear

Programming

- Specifically, consider **subdividing** the path after one step
- Then an optimal path from s consists of:
 - An **optimal first action** a^*
 - Followed by **an optimal path** from $s' = \text{succ}(s, a^*)$
- Therefore the value of an optimal path must satisfy:

$$V^*(s) = \max_{a \in \mathcal{A}} R(s, a) + V^*(s')$$



Deterministic Value Iteration

Marcello
Restelli

Policy Search

Dynamic
Programming

Policy Iteration

Value Iteration

Extensions to

Dynamic

Programming

Linear

Programming

- If we know the solution to **subproblems** $V^*(s')$
- Then it is easy to construct the solution to $V^*(s)$

$$V^*(s) \leftarrow \max_{a \in \mathcal{A}} R(s, a) + V^*(s')$$

- The idea of value iteration is to apply these updates **iteratively**
- e.g., Starting from the goal and **working backward**



Value Iteration in MDPs

Marcello
Restelli

Policy Search

Dynamic
Programming

Policy Iteration

Value Iteration

Extensions to

Dynamic

Programming

Linear
Programming

- Many MDPs **don't have a finite horizon**
- They are typically **loopy**
- So there is **no "end"** to work backwards from
- However, we can still **propagate information backwards**
- Using **Bellman optimality equation** to backup $V(s)$ from $V(s')$
- Each subproblem is "easier" due to **discount factor** γ
- Iterate until **convergence**



Principle of Optimality in MDPs

Marcello
Restelli

Policy Search

Dynamic
Programming

Policy Iteration

Value Iteration

Extensions to

Dynamic

Programming

Linear

Programming

Theorem

A policy $\pi(a|s)$ achieves the **optimal value** from state s , $V^\pi(s) = V^*(s)$, if and only if

- For any state s' reachable from s
- π achieves the optimal value from state s' , $V^\pi(s') = V^*(s')$

So an optimal policy $\pi^*(a|s)$ must consist of:

- an optimal first action a^*
- followed by an optimal policy from successor state s'

$$V^*(s) = \max_{a \in \mathcal{A}} R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V^*(s')$$



Value Iteration

Marcello
Restelli

Policy Search

Dynamic
Programming

Policy Iteration

Value Iteration

Extensions to

Dynamic
Programming

Linear
Programming

- **Problem:** find optimal policy π
- **Solution:** iterative application of Bellman optimality backup
- $V_1 \rightarrow V_2 \rightarrow \dots \rightarrow V^*$
- Using **synchronous backups**
 - At each iteration $k + 1$
 - For all states $s \in \mathcal{S}$
 - Update $V_{k+1}(s)$ from $V_k(s')$
- Unlike policy iteration there is **no explicit policy**
- **Intermediate** value functions **may not correspond** to any policy

Value Iteration demo:

<http://www.cs.ubc.ca/~poole/demos/mdp/vi.html>



Convergence and Contractions

Marcello
Restelli

Policy Search

Dynamic
Programming

Policy Iteration
Value Iteration

Extensions to
Dynamic
Programming

Linear
Programming

Define the max-norm: $\|V\|_{\infty} = \max_s |V(s)|$

Theorem

Value Iteration converges to the optimal state-value function

$$\lim_{k \rightarrow \infty} V_k = V^*$$

Proof.

$$\|V_{k+1} - V^*\|_{\infty} = \|T^* V_k - T^* V^*\|_{\infty} \leq \gamma \|V_k - V^*\|_{\infty} \leq \dots \leq \gamma^{k+1} \|V_0 - V^*\|_{\infty} \rightarrow 0 \quad \square$$

Theorem

$$\|V_{i+1} - V_i\|_{\infty} < \epsilon \Rightarrow \|V_{i+1} - V^*\|_{\infty} < \frac{2\epsilon\gamma}{1-\gamma}$$



Synchronous Dynamic Programming Algorithms

Marcello
Restelli

Policy Search

Dynamic
Programming

Policy Iteration

Value Iteration

Extensions to
Dynamic
Programming

Linear
Programming

Problem	Bellman Equation	Algorithm
Prediction	Bellman Expectation Equation	Policy Evaluation (Iterative)
Control	Bellman Expectation + Greedy Policy Improvement	Policy Iteration
Control	Bellman Optimality Equation	Value Iteration

- Algorithms are based on **state-value function** $V^\pi(s)$ or $V^*(s)$
- Complexity $O(mn^2)$ **per iteration**, for m actions and n states
- Could also apply to **action-value function** $Q^\pi(s, a)$ or $Q^*(s, a)$
- Complexity $O(m^2n^2)$ **per iteration**



Efficiency of DP

Marcello
Restelli

Policy Search

Dynamic
Programming

Policy Iteration

Value Iteration

Extensions to

Dynamic

Programming

Linear
Programming

- To find optimal policy is **polynomial** in the number of states...
- **but**, the number of states is often astronomical, e.g., often growing **exponentially** with the number of state variables: **curse of dimensionality**
- In practice, classical DP can be applied to problems with a few millions states
- **Asynchronous DP** can be applied to larger problems, and appropriate for parallel computation
- It is surprisingly **easy** to come up with MDPs for which methods are not practical



Complexity of DP

Marcello
Restelli

Policy Search

Dynamic
Programming

Policy Iteration
Value Iteration

Extensions to
Dynamic
Programming

Linear
Programming

- DP methods are **polynomial time** algorithms for **fixed-discounted** MDPs
- **Value Iteration:** $O(|\mathcal{S}|^2 |\mathcal{A}|)$ for each iteration
- **Policy Iteration:** Cost of policy evaluation + Cost of policy iteration
 - Policy evaluation:
 - Linear system of equations: $O(|\mathcal{S}|^3)$ or $O(|\mathcal{S}|^{2.373})$
 - Iterative: $O\left(|\mathcal{S}|^2 \frac{\log(\frac{1}{\epsilon})}{\log(\frac{1}{\gamma})}\right)$
 - Policy improvement: recently proven to be $O\left(\frac{|\mathcal{A}|}{1-\gamma} \log\left(\frac{|\mathcal{S}|}{1-\gamma}\right)\right)$
- **Each iteration** of PI is computationally **more expensive** than each iteration of VI
- PI typically requires fewer iterations to converge than VI
- **Exponentially faster** than any **direct policy search**
- Number of states often **grows exponentially** with the number of state variables



Asynchronous Dynamic Programming

Marcello
Restelli

Policy Search

Dynamic
Programming

Policy Iteration

Value Iteration

Extensions to

Dynamic
Programming

Linear

Programming

- DP methods described so far used **synchronous** backups
 - i.e., all state are backed up in **parallel**
- **Asynchronous** DP backs up states **individually**, in any order
- For each **selected state**, apply the appropriate backup
- Can significantly **reduce computation**
- Guaranteed to **converge** if all states continue to be selected
- Three ideas for asynchronous DP:
 - In-place DP
 - Prioritized sweeping
 - Real-time DP



In-place Dynamic Programming

Marcello
Restelli

Policy Search

Dynamic
Programming

Policy Iteration

Value Iteration

Extensions to

Dynamic
Programming

Linear

Programming

- **Synchronous** value iteration stores **two copies** of value function

for all $s \in \mathcal{S}$

$$V_{new}(s) \leftarrow \max_{a \in \mathcal{A}} \left(R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V_{old}(s') \right)$$

$$V_{old} \leftarrow V_{new}$$

- **In-place** value iteration only stores **one copy** of value function

for all $s \in \mathcal{S}$

$$V(s) \leftarrow \max_{a \in \mathcal{A}} \left(R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V(s') \right)$$



Prioritized Sweeping

Marcello
Restelli

Policy Search

Dynamic
Programming

Policy Iteration
Value Iteration

Extensions to
Dynamic
Programming

Linear
Programming

- Use the magnitude of **Bellman error** to guide state selection, e.g.,

$$\left| \max_{a \in \mathcal{A}} \left(R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V_k(s') \right) - V(s) \right|$$

- Backup the state with the **largest** remaining Bellman error
- **Update** Bellman error of affected states after each backup
- Requires knowledge of **reverse dynamics** (predecessor states)
- Can be implemented efficiently by maintaining a **priority queue**



Real-Time Dynamic Programming

Marcello
Restelli

Policy Search

Dynamic
Programming

Policy Iteration

Value Iteration

Extensions to

Dynamic
Programming

Linear
Programming

- **Idea**: only states that are **relevant** to agent
- Use agent's **experience** to guide the **selection** of states
- After each time-step s_t, a_t, r_{t+1}

$$a_t \in \arg \max_{a \in \mathcal{A}} \left(R(s_t, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s_t, a) V(s') \right)$$

- Backup the state s_t

$$V(s_t) \leftarrow \max_{a \in \mathcal{A}} \left(R(s_t, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s_t, a) V(s') \right)$$

Theorem

If $V_0 \geq V^$ then $\exists \bar{t}$ such that a_t are optimal for all $t \geq \bar{t}$ (where $\bar{t} < \infty$ with probability 1)*



Full-Width Backups

Marcello
Restelli

Policy Search

Dynamic
Programming

Policy Iteration
Value Iteration

Extensions to
Dynamic
Programming

Linear
Programming

- Dynamic programming uses **full-width** backups
- For each backup (synchronous or asynchronous)
 - **Every** successor state and action is considered
 - Using knowledge of the MDP **transitions** and **reward function**
- Dynamic programming is effective for **medium-size** problems (millions of states)
- For large problems dynamic programming suffers Bellman's **curse of dimensionality**
 - Number of states $n = |\mathcal{S}|$ grows **exponentially** with number of states variables
- Even **one backup** can be too **expensive**



Sample Backups

Marcello
Restelli

Policy Search

Dynamic
Programming

Policy Iteration

Value Iteration

Extensions to

Dynamic

Programming

Linear
Programming

- Reinforcement Learning techniques exploit **sample backups**
- Sample backups do not use reward function R and transition dynamics P
- Uses sample rewards and sample transitions $\langle s, a, s', r \rangle$
- Advantages
 - **Model-free**: no prior knowledge of MDP required
 - Breaks the curse of dimensionality through sampling
 - Cost of backups is constant, independent of $n = |\mathcal{S}|$



Approximate Dynamic Programming

Marcello
Restelli

Policy Search

Dynamic
Programming

Policy Iteration
Value Iteration

Extensions to
Dynamic
Programming

Linear
Programming

- **Approximate** the value function
- Using a **function approximator** $V^\theta(s) = f(s, \theta)$
- Apply dynamic programming to V^θ
 - e.g., **Fitted Value Iteration** repeats at each iteration k
 - Sample states $\tilde{\mathcal{S}} \subseteq \mathcal{S}$
 - For each state $s \in \tilde{\mathcal{S}}$, estimate target value using Bellman optimality equation

$$\tilde{V}_k(s) = \max_{a \in \mathcal{A}} \left(R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V_k^\theta(s') \right)$$

- Train next value function V_{k+1}^θ using targets $\{\langle s, \tilde{V}_k(s) \rangle\}$



Infinite Horizon Linear Programming

Marcello
Restelli

Policy Search

Dynamic
Programming

Policy Iteration

Value Iteration

Extensions to

Dynamic
Programming

Linear
Programming

- Recall, at value iteration convergence we have

$$\forall s \in \mathcal{S} : \quad V^*(s) = \max_{a \in \mathcal{A}} \left\{ R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V^*(s') \right\}$$

- LP formulation to find V^* :

$$\begin{array}{ll} \min_V & \sum_{s \in \mathcal{S}} \mu(s) V(s) \\ \text{s. t.} & V(s) \geq R(s, a) + \sum_{s' \in \mathcal{S}} P(s'|s, a) V(s'), \quad \forall s \in \mathcal{S}, \forall a \in \mathcal{A} \end{array}$$

- $|\mathcal{S}|$ variables
- $|\mathcal{S}||\mathcal{A}|$ constraints

Theorem

V^ is the solution of the above LP.*



Theorem Proof

Marcello
Restelli

Policy Search

Dynamic
Programming

Policy Iteration

Value Iteration

Extensions to

Dynamic
Programming

Linear
Programming

Let T^* be the **optimal Bellman operator**, then the LP can be written as:

$$\begin{array}{ll} \min_V & \mu^\top V \\ \text{s. t.} & V \geq T^*(V) \end{array}$$

- **Monotonicity property:** if $U \geq V$ then $T^*(U) \geq T^*(V)$.
- Hence, if $V \geq T^*(V)$ then $T^*(V) \geq T^*(T^*(V))$, and by **repeated application**,
$$V \geq T^*(V) \geq T^{*2}(V) \geq T^{*3}(V) \geq \dots \geq T^{*\infty}(V) = V^*$$
- Any **feasible solution** to the LP must satisfy $V \geq T^*(V)$, and hence must satisfy $V \geq V^*$
- Hence, assuming all entries μ are positive, V^* is the **optimal solution** to the LP



Dual Linear Program

Marcello
Restelli

Policy Search

Dynamic
Programming

Policy Iteration

Value Iteration

Extensions to

Dynamic

Programming

Linear
Programming

$$\begin{aligned} \max_{\lambda} \quad & \sum_{s \in \mathcal{S}} \sum_{a \in \mathcal{A}} \lambda(s, a) R(s, a) \\ \text{s. t.} \quad & \sum_{a' \in \mathcal{A}} \lambda(s', a') = \mu(s) + \gamma \sum_{s \in \mathcal{S}} \sum_{a \in \mathcal{A}} \lambda(s, a) P(s' | s, a), \quad \forall s' \in \mathcal{S} \\ & \lambda(s, a) \geq 0, \quad \forall s \in \mathcal{S}, a \in \mathcal{A} \end{aligned}$$

- **Interpretation**

- $\lambda(s, a) = \sum_{t=0}^{\infty} \gamma^t \mathbb{P}(s_t = s, a_t = a)$
- Equation 2: ensures λ has the above meaning
- Equation 1: maximize expected discounted sum of rewards

- **Optimal policy:** $\pi^*(s) = \arg \max_a \lambda(s, a)$



Complexity of LP

Marcello
Restelli

Policy Search

Dynamic
Programming

Policy Iteration

Value Iteration

Extensions to

Dynamic
Programming

Linear
Programming

- LP **worst-case** convergence guarantees are better than those of DP methods
- LP methods become **impractical** at a much smaller number of states than DP methods do